

AWS Solution Architect Associate

Version : C03

Domain 2

Task 1

**Design scalable and loosely coupled
architectures**

REST API

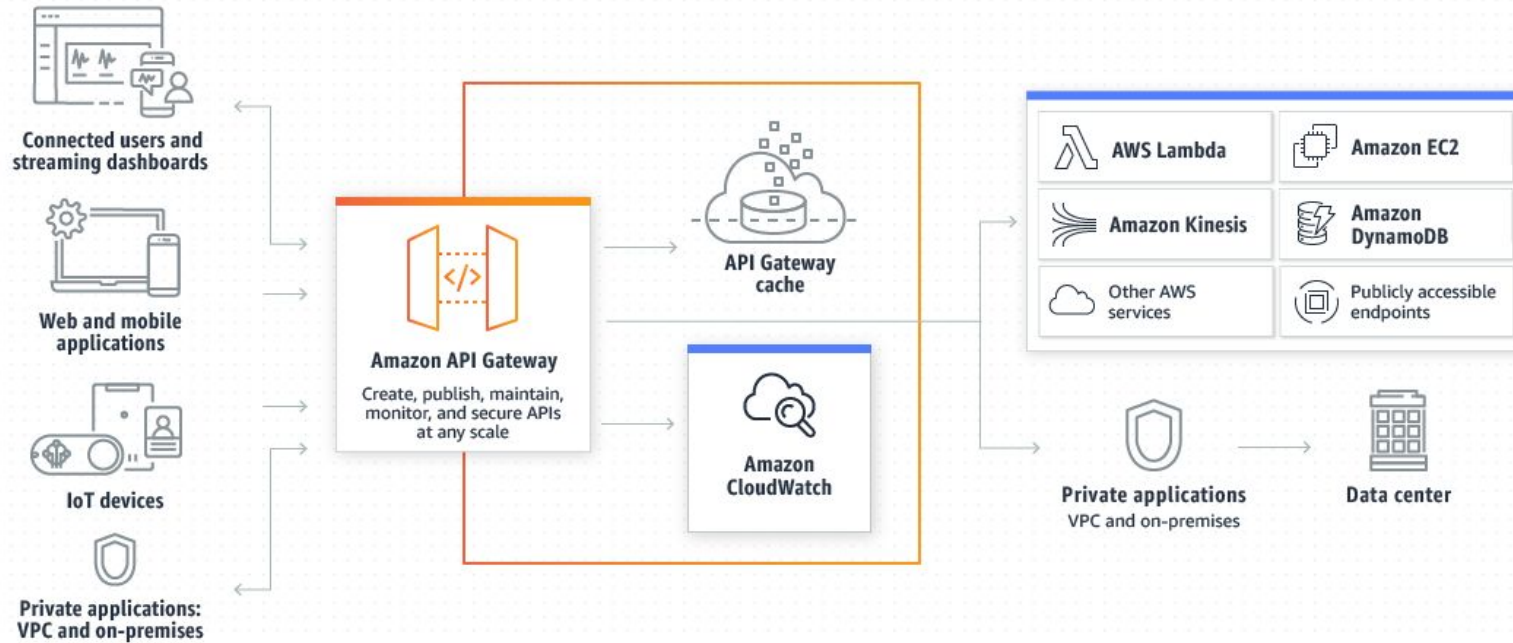
- AWS REST API is a scalable, secure and easy-to-use interface for accessing AWS resources.
 - It allows applications to access and control the AWS services using HTTP requests.
 - Provides a set of standard endpoints and methods for accessing various AWS services
 - Tools like cURL or SDKs can be used to make requests to the API
 - It supports features like authentication, authorization and pagination to ensure secure and efficient access to AWS resources.
 - REST APIs and HTTP APIs are both RESTful API products.
 - ◆ REST APIs support more features than HTTP APIs
 - ◆ HTTP APIs are offered at a lower price.
 - ◆ REST APIs features - API keys, per-client throttling, request validation, AWS WAF integration or private API endpoints
-

Amazon API Gateway




- Amazon API Gateway is an AWS service for creating, publishing, maintaining, monitoring and securing REST, HTTP, and WebSocket APIs at any scale.
 - APIs can be created that access AWS or other web services, as well as data stored in the AWS Cloud.
 - API Gateway creates RESTful APIs that :
 - ◆ Are HTTP-based
 - ◆ Enable stateless client-server communication
 - ◆ Implement standard HTTP methods such as GET, POST, PUT, PATCH and DELETE
 - API Gateway creates WebSocket APIs that :
 - ◆ Adhere to the WebSocket protocol which enables stateful, full-duplex communication between client and server.
 - ◆ Route incoming messages based on message content.
-

Architecture of Amazon API Gateway

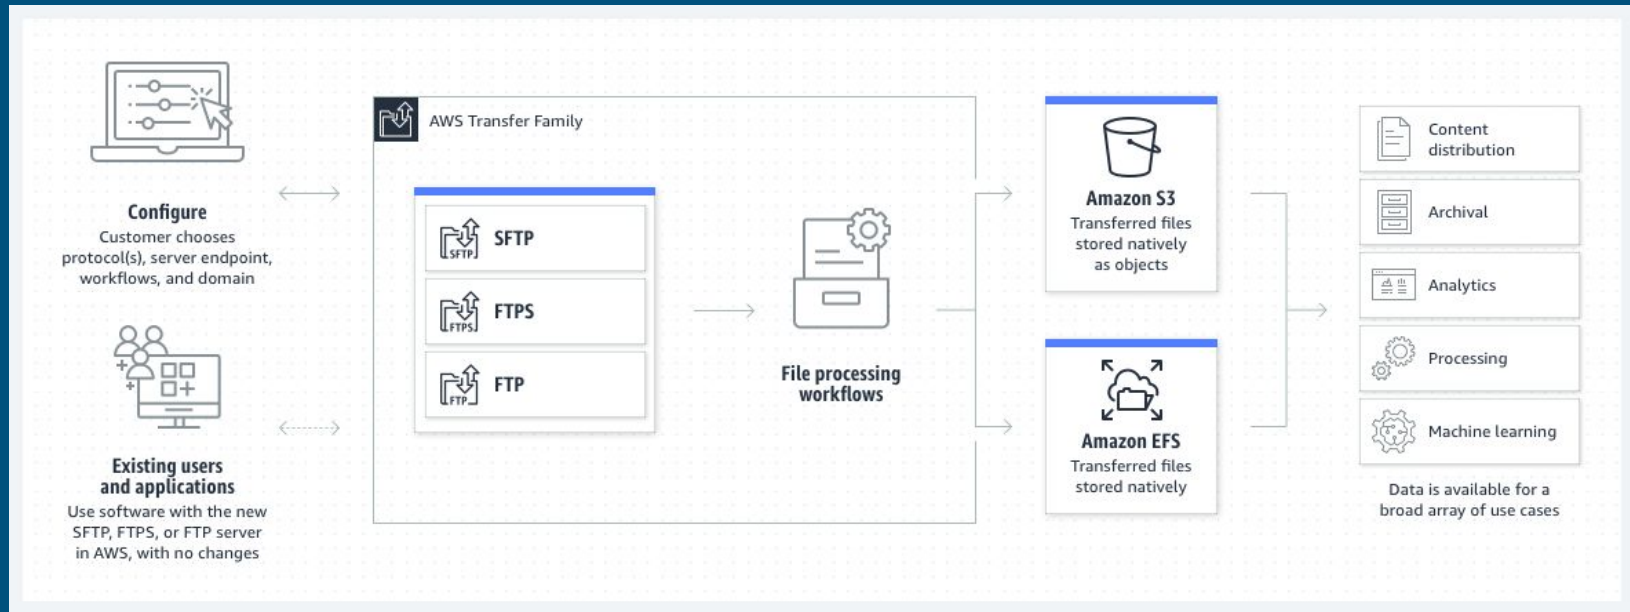


AWS Transfer Family



- AWS Transfer Family is a secure transfer service that enables you to transfer files into and out of AWS storage services.
 - It is a part of the AWS Cloud platform.
 - Support over following protocols
 - ◆ Secure File Transfer Protocol (SFTP-v3)
 - ◆ File Transfer Protocol Secure (FTPS)
 - ◆ File Transfer Protocol (FTP)
 - ◆ Applicability Statement (AS2)
 - File transfer workflows can seamlessly migrate, automate and monitor the file transfer workflows
 - Supports transferring data from or to the following Amazon web services
 - ◆ Amazon S3
 - ◆ Amazon Elastic File System
 - Use Cases
 - ◆ Modernize your managed file transfers
 - ◆ Gain insights by growing your data lake
 - ◆ Improve collaboration across your trading partner network
 - ◆ Expand your content distribution business
-

AWS Transfer Family



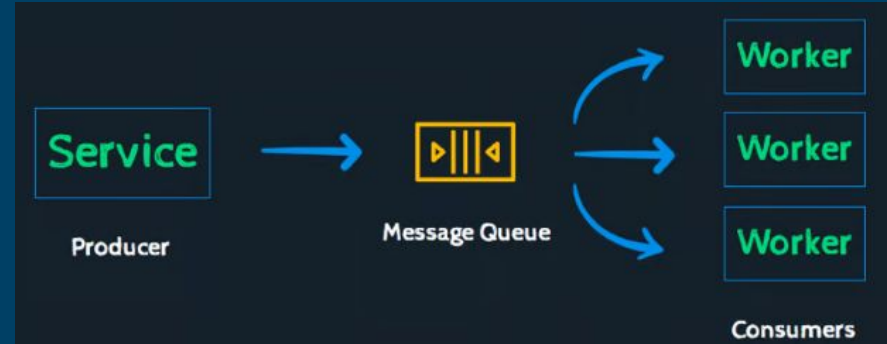
Amazon Simple Queue Service



- Amazon SQS aka Amazon Simple Queue Service
- Offers a secure, durable and available hosted queue that integrate and decouple distributed software systems and components.
- It offers common constructs such as dead-letter queues and cost allocation tags.
- Benefits of using SQS
 - ◆ Security
 - ◆ Durability
 - ◆ Availability
 - ◆ Scalability
 - ◆ Reliability



Amazon Simple Queue Service



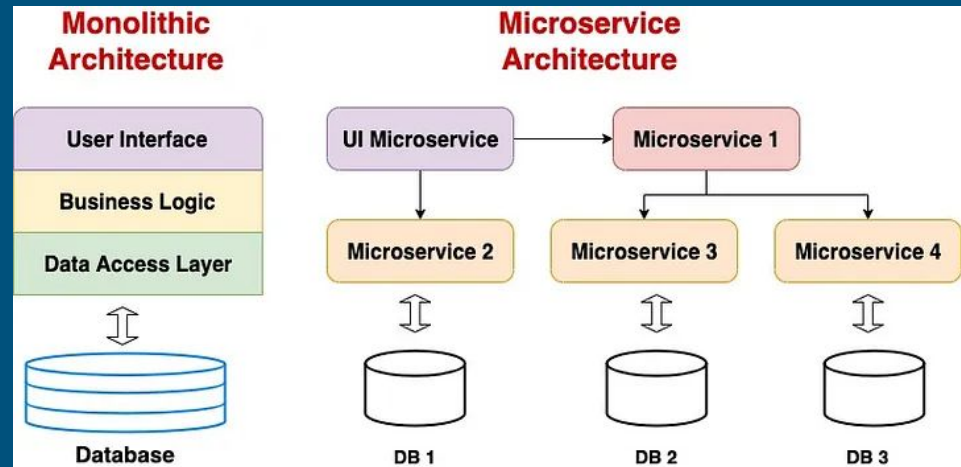
- Queues for Message Storage
- Producer-Consumer Model
- Message Lifecycle
- Dead-Letter Queues (DLQs)
- Scaling and Fault Tolerance

Caching Strategies

- Refers to the methods for caching data in Amazon's cloud-based infrastructure to maximize its performance and availability.
 - AWS provides several caching options
 - ◆ Amazon ElastiCache
 - ◆ Amazon S3 Caching Acceleration
 - ◆ AWS Lambda Caching
 - ◆ Amazon File Cache
 - ◆ Amazon DynamoDB Accelerator (DAX)
 - ◆ Amazon Greengrass
 - Each strategy has unique performance characteristics and is suited for different workloads.
 - The choice of caching strategy depends on the specific requirements of the application including data access frequency, access patterns and data size.
 - Caching strategies can help to reduce latency, improve throughput and reduce costs associated with storage and data retrieval.
 - Different caching strategies
 - ◆ **Lazy loading**
 - ◆ **Write-Through**
-

Design Principles for Microservices

- Approach of breaking down a large monolithic application into smaller, independent services that are more scalable, flexible and maintainable
- Key principles
 - ◆ **Loose coupling**
 - ◆ **High cohesion**
 - ◆ **Service contract design**
 - ◆ **Failure tolerance**
 - ◆ **Independent deployment**



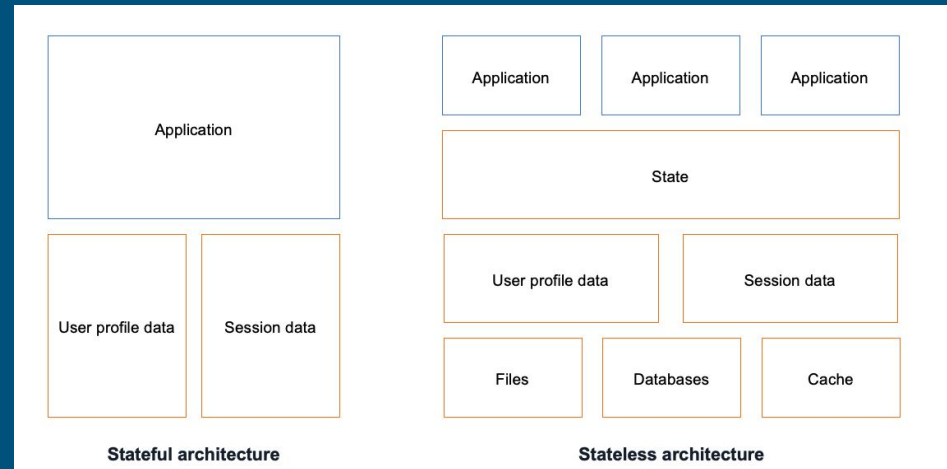
Stateless/Stateful Workloads

→ Stateful applications

- ◆ Simple to deploy
- ◆ Saves client session data on the server, allowing for faster processing and improved performance.
- ◆ Stateful applications excel in predictable workloads and offer consistent user experiences.

→ Stateless architectures

- ◆ Align with the demands of dynamic workload and changing business requirements.
- ◆ Stateless application design can increase flexibility with horizontal scaling and dynamic deployment.
- ◆ This flexibility helps applications handle sudden spikes in traffic, maintain resilience to failures and optimize cost.



Event Driven Architectures


- An event-driven architecture uses events to trigger and communicate between decoupled services and is common in modern applications built with microservices.
 - An event is a change in state or an update, like an item being placed in a shopping cart on an e-commerce website.
 - Events can either carry the state or events can be identifiers
 - Key components :
 - ◆ **Event producers**
 - ◆ **Event routers**
 - ◆ **Event consumers**
 - A producer publishes an event to the router, which filters and pushes the events to consumers.
 - Producer services and consumer services are decoupled, which allows them to be scaled, updated and deployed independently.
-

Horizontal/Vertical Scaling


- AWS horizontal scaling and vertical scaling are two different approaches to scaling applications in the cloud.
 - **Horizontal scaling**
 - ◆ Involves adding more instances of a running application to handle increased traffic or workload.
 - ◆ This can be done using Amazon Elastic Load Balancing (ELB) and Amazon Auto Scaling.
 - **Vertical scaling**
 - ◆ Involves increasing the resources available to a single instance, such as CPU, memory, or storage.
 - ◆ This can be done using Amazon EC2 instances with larger instance types.
 - Both approaches have advantages and disadvantages and the choice of which to use depends on the specific needs of the application.
-

Horizontal/Vertical Scaling

Horizontal Scaling



Vertical Scaling



Content Delivery Network

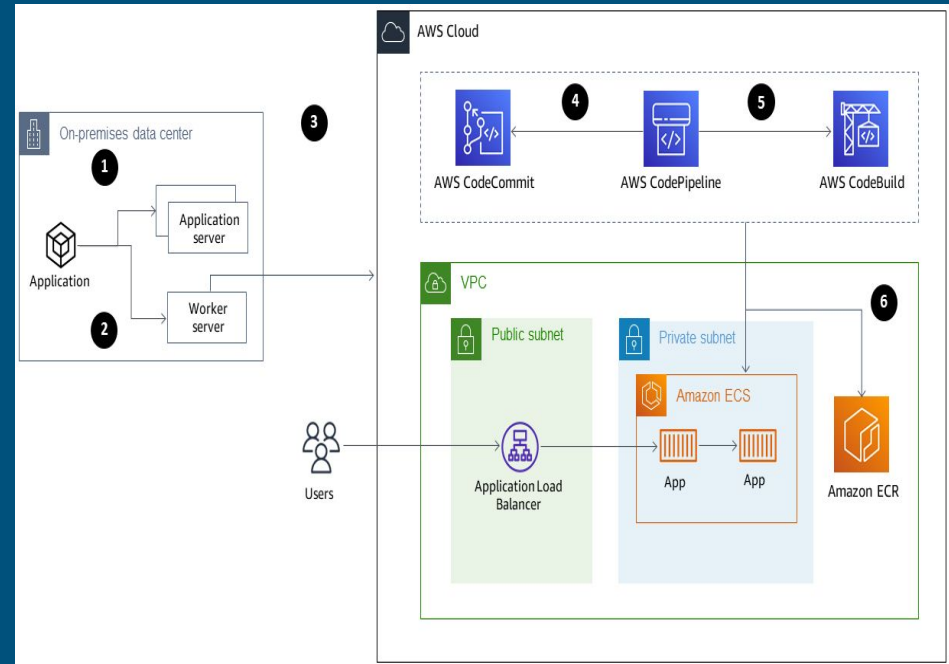
- CDN is a network of interconnected servers that speeds up webpage loading for data-heavy applications.
 - It can stand for content delivery network or content distribution network.
 - Use case : When a user visits a website, data from that website's server has to travel across the internet to reach the user's computer. If the user is located far from that server, it will take a long time to load a large file, such as a video or website image. Instead, the website content is stored on CDN servers geographically closer to the users and reaches their computers much faster.
 - The primary purpose is to reduce latency/delay in communication created by a network's design
 - Benefits of CDN
 - ◆ Reduce Page load time
 - ◆ Reduce Bandwidth costs
 - ◆ Increase Content availability
 - ◆ Improve Website security
 - CDN can deliver 2 types of content
 - ◆ Static Content
 - ◆ Dynamic Content
-

Content Delivery Network (Edge Accelerators)

- AWS Edge Accelerator is a content delivery network (CDN) service that helps deliver content faster to end-users by caching content closer to them.
 - It is designed to improve the performance of web applications and APIs by reducing latency and improving throughput.
 - AWS Edge Accelerator can be used to cache static content, such as images and videos, as well as dynamic content, such as web pages and APIs.
 - Supports a variety of protocols - HTTP, HTTPS, and WebSocket and can be integrated with other AWS services such as Amazon CloudFront and Amazon API Gateway.
-

Migrating Applications to containers

- AWS offers several services for migrating applications to containers
 - ◆ **AWS Elastic Container Service**
 - ◆ **AWS Elastic Kubernetes Service**
 - ◆ **AWS Fargate**
- These services provide a scalable and flexible infrastructure for deploying and managing containerized applications.
- It offers tools and resources for containerizing applications
 - ◆ **AWS Container Registry**
 - ◆ **AWS CodeBuild**
- AWS provides guidance and best practices for migrating applications to containers, including considerations for security, performance and scalability.



Migrating Applications to containers

→ Steps for migrating applications to containers using AWS :

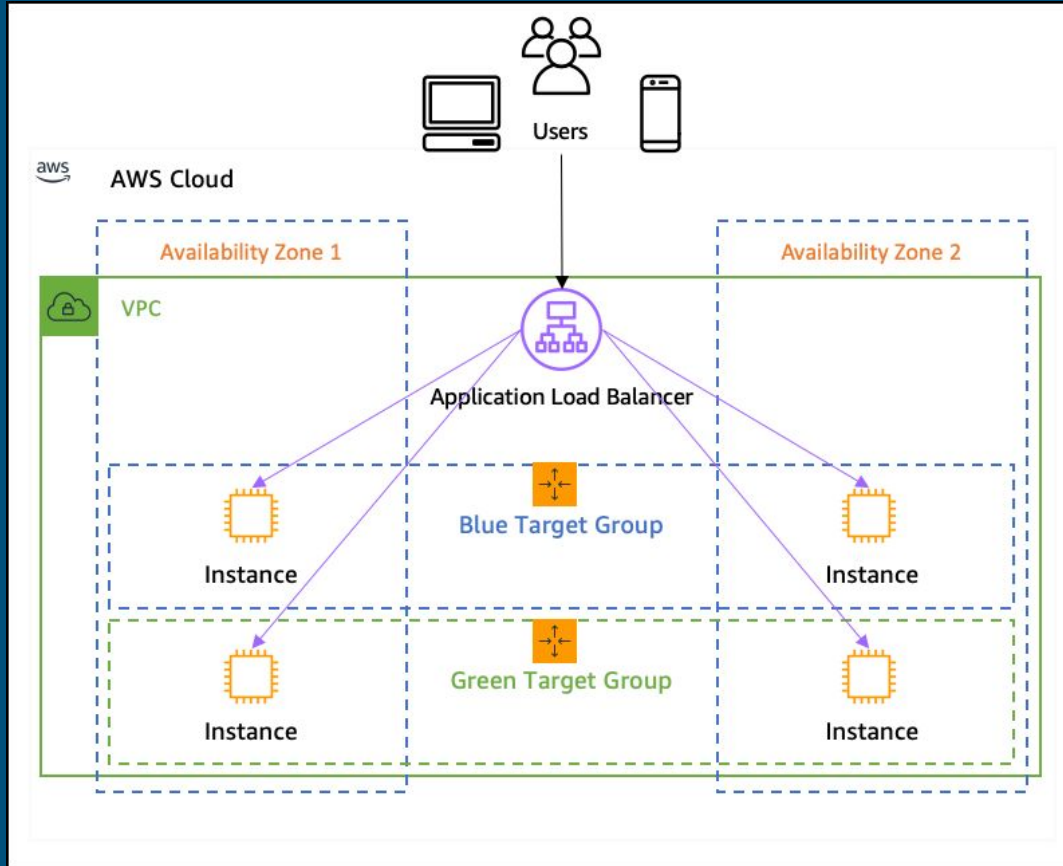
- ◆ **Assess your application** - Determine which containers are appropriate for your application and how they should be configured.
 - ◆ **Modify your application** - Modify your application code to run in containers.
 - ◆ **Choose a container service** - Select the container service that best meets your needs.
 - ◆ **Create container images** - Create container images for your application using a tool such as Docker.
 - ◆ **Deploy to the container service** - Deploy your containers
-

Load Balancing



- Refers to the process of distributing network traffic across multiple servers or instances to ensure that no single server becomes overloaded.
 - This helps to improve the performance, scalability, and reliability of cloud-based applications and services.
 - AWS Load Balancing
 - ◆ Service offered by (AWS) that helps distribute incoming network traffic across multiple instances, containers or IP addresses
 - ◆ Improves application performance and availability, simplify scaling and configuration and enhance security.
 - ◆ Types of load balancing offered by AWS :
 - **Application Load Balancers**
 - **Network Load Balancers**
 - **Load Balancer Target Groups**
 - **EC2 Auto Scaling Groups**
-

Load Balancing



Multi-Tier Architecture

- AWS multi-tier architecture is a **type of system design** that involves breaking down an application into smaller, more manageable components, which are distributed across multiple layers.
 - This provides better scalability, performance and cost efficiency.
 - The layers includes
 - ◆ **Presentation**
 - ◆ **Application**
 - ◆ **Business logic**
 - ◆ **Data storage and retrieval**
 - ◆ **Data processing**
 - This reduces coupling and makes the system easier to manage and maintain.
 - Different AWS tools/services to implement a multi-tier architecture -
 - ◆ **Elastic Load Balancing**
 - ◆ **AWS Auto Scaling**
 - ◆ **Amazon Elastic Container Service.**
-

Serverless technologies

- In the serverless model, a cloud provider manages the provisioning, scaling, and maintaining the underlying infrastructure.
 - The cloud provider handles tasks like operating system management, security patches, file system and capacity management, load balancing, monitoring and logging.
 - Serverless use cases :
 - ◆ Business process automation
 - ◆ Real-time data analytics
 - ◆ Batch processing
 - ◆ Stateless application development
 - Types of serverless architecture
 - ◆ **Functions as a service** - Can be used to write and deploy the function code directly to the cloud infrastructure
 - ◆ **Backend as a service** - Provides access to backend functions using an API
-

Serverless technologies (AWS Services)

→ **Compute**

- ◆ AWS Lambda
- ◆ AWS Fargate

→ **Application Integration**

- ◆ Amazon EventBridge
- ◆ AWS Steps Functions
- ◆ Amazon SQS
- ◆ Amazon SNS
- ◆ Amazon API Gateway
- ◆ AWS AppSync

→ **Data Store**

- ◆ Amazon S3
 - ◆ Amazon EFS
 - ◆ Amazon DynamoDB
 - ◆ Amazon RDS Proxy
 - ◆ Amazon Aurora Serverless
 - ◆ Amazon Redshift Serverless
 - ◆ Amazon Neptune Serverless
 - ◆ Amazon OpenSearch Serverless
 - ◆ Amazon ElastiCache Serverless
-

Storage types

→ Three main cloud storage types and the corresponding AWS services :

◆ **Object Storage**

- Amazon S3
- Amazon Glacier
- Amazon S3 Glacier Deep Archive
- Amazon S3 Intelligent-Tiering

◆ **File Storage**

- Amazon EFS
- Amazon FSx for Windows File Server
- Amazon FSx for Lustre
- AWS Storage Gateway

◆ **Block Storage**

- Amazon EBS
 - Amazon EC2 instance store
 - Amazon EFS (NFSv4)
-

Container Orchestration

- Involves managing , automating the deployment, scaling and operation of containerized applications.
 - AWS service for container orchestration
 - ◆ **Amazon Elastic Container Service**
 - ◆ **Amazon Elastic Kubernetes Service**
 - ◆ **AWS Fargate**
 - ◆ **AWS App Runner**
 - **Amazon Elastic Container Service -**
 - ◆ Fully managed container orchestration service
 - ◆ Supports Docker containers.
 - ◆ Allows to run\manage containerized applications
 - ◆ Task scheduling, service discovery, auto scaling and integration with other AWS services.
 - **Amazon Elastic Kubernetes Service**
 - ◆ Fully managed Kubernetes service
 - ◆ Fully managed control plane and integration with AWS services
 - ◆ Simplifies the deployment, management, and scaling of Kubernetes applications
 - ◆ Can run containerized applications using standard Kubernetes APIs and tooling.
-

Container Orchestration

- **AWS Fargate**
 - ◆ Serverless compute engine for containers
 - ◆ Container requirements such as CPU and memory can be specified and AWS takes care of provisioning and scaling the compute resources automatically.
 - ◆ Integrates seamlessly with Amazon ECS and Amazon EKS
 - ◆ No need to manage servers or clusters
 - **AWS App Runner**
 - ◆ Fully managed service that deploy and run containerized web applications at scale.
 - ◆ Automatically builds and deploys the container images from source code or Dockerfiles, manages the infrastructure, and scales the application based on traffic.
 - ◆ Suitable for running web applications, APIs, and microservices with minimal management overhead.
-

Container Orchestration

AWS container options by layer



AWS
App Runner



AWS Elastic
Beanstalk



Amazon
Lightsail

Provisioning



Amazon
ECS



Amazon
EKS



ROSA

Orchestration



Amazon
EC2

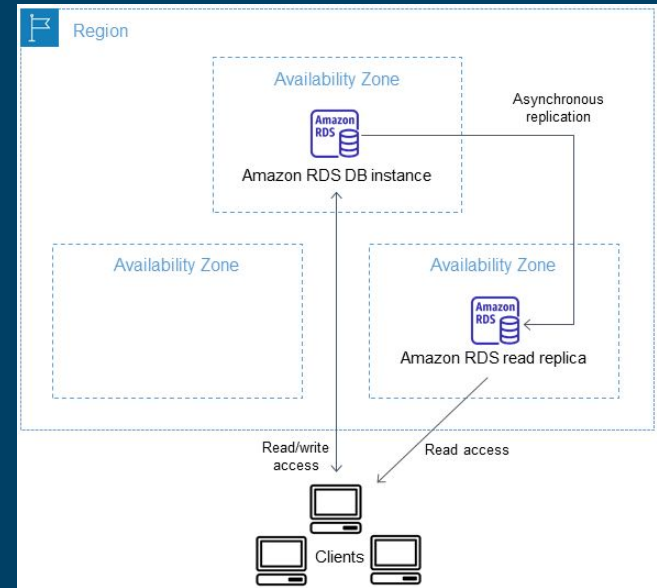


AWS
Fargate

Capacity

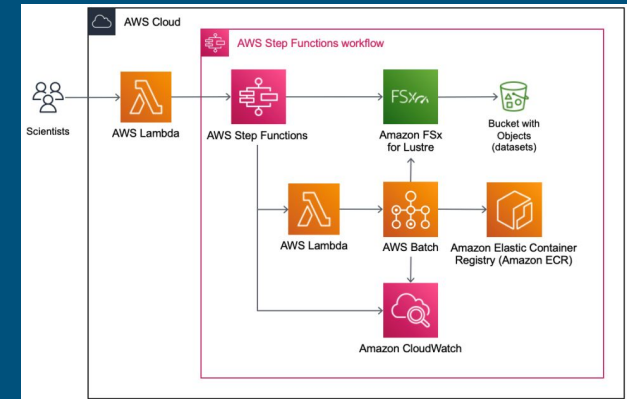
Read Replicas

- A read replica is a read-only copy of a DB instance
- Load on the primary DB instance can be reduced by routing queries from your applications to the read replica
- Can elastically scale out beyond the capacity constraints of a single DB instance for read-heavy database workloads



Workflow Orchestration

- Technology that allows you to create and manage workflows to automate tasks and processes.
- Workflow orchestration is a key component of AWS step functions
- **AWS Step Functions**
 - ◆ A serverless orchestration service
 - ◆ Combines AWS Lambda functions and other AWS services to build to scalable, distributed applications using state machines
 - ◆ Based on state machine and tasks
 - ◆ State machine is a workflow
 - ◆ Task is a state in workflow that represents a single unit of work that another AWS service performs
 - ◆ Scales horizontally and provides fault-tolerant workflows
 - ◆ Workflow types
 - **Standard Workflows**
 - Exactly once workflow transformation
 - **Express Workflows**
 - At-Least once workflow transformation



AWS Solution Architect Associate

Version : C03

Domain 2

Task 1

The END